

Homework #5 Report: Airflow DAGs with Snowflake Integration

Abstract

This report documents the implementation of Homework #5. Apache Airflow orchestrates an ETL and a forecasting pipeline with Snowflake as the data warehouse. Screenshots demonstrate configuration and successful runs. Concise, relevant code snippets are embedded per the assignment requirement.

I. INTRODUCTION

This submission shows:

- Airflow DAGs (`yfinance_etl`, `ml_forecast_tf`, and a simplified assignment DAG).
- Proper use of Airflow **Connections** and **Variables**.
- SQL **transaction** semantics for full refresh.
- Evidence of successful runs via the Airflow UI logs.

II. AIRFLOW WEB UI: DAGS LOADED

Figure 1 confirms four DAGs are discovered by the scheduler and are active.

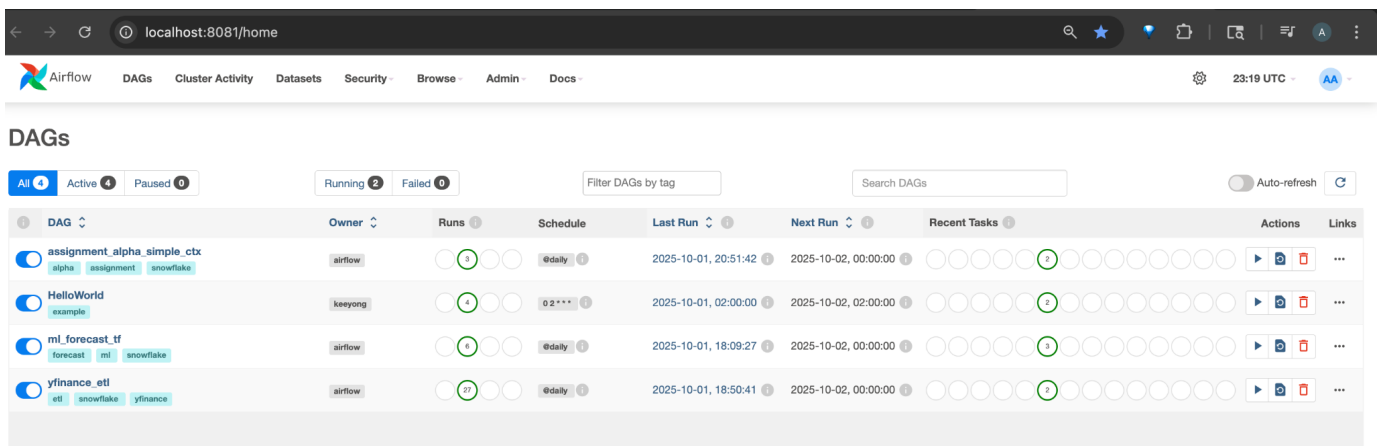


Fig. 1: Airflow Web UI showing all DAGs loaded.

III. ADMIN → CONNECTION (SNOWFLAKE)

The connection `snowflake_catfish` was created with type *Snowflake*. Account, warehouse, database, schema, and role are set in the Extra/fields as required.

The screenshot shows the 'Edit Connection' form in the Airflow web interface. At the top, there's a navigation bar with links for DAGs, Cluster Activity, Datasets, Security, Browse, Admin, and Docs. Below the navigation bar is a warning message: 'Warning: Fields that are currently populated can be modified but cannot be deleted. To delete data from a field, delete the Connection object and create a new one.' The form itself has a title 'Edit Connection' and contains several fields: 'Connection Id' with the value 'snowflake_catfish', 'Connection Type' with the value 'Snowflake' and a note 'Connection Type missing? Make sure you've installed the corresponding Airflow Provider Package.', 'Description' (empty), 'Schema' with the value 'RAW', 'Login' with the value 'CATFISH', and 'Password' with the value 'snowflake password'.

Fig. 2: Airflow Connection configuration for Snowflake.

How to Configure the Airflow Snowflake Connection (UI)

Path: Admin → Connections → + Add a new record

- 1) **Connection Id:** snowflake_catfish (must match the code)
- 2) **Connection Type:** Snowflake
- 3) **Login:** Snowflake user, e.g., CATFISH
- 4) **Password:** Snowflake password for the above user
- 5) **Schema:** Default schema for this connection, e.g., RAW
- 6) **Extra (JSON):** Provide account/warehouse/database/role (and optionally schema):

```

1  {
2    "account": "xxxx-xxxx",
3    "warehouse": "COMPUTE_WH",
4    "database": "YOUR_DB",
5    "schema": "RAW",
6    "role": "SYSADMIN"
7  }

```

Why this matters.

The DAG uses `BaseHook.get_connection("snowflake_catfish")` to read these values at runtime. Keeping these secrets in the Airflow UI avoids hardcoding and makes deployments portable across environments.

Where it is used in code (exact snippet).

How the Snowflake Connection is Used in Assignment_5.py

In the DAG file (Assignment_5.py), the Snowflake connection is referenced through an Airflow Variable, which points to the connection ID defined in the Admin → Connections page.

Code Snippet (from Assignment_5.py).

```

1  from airflow.providers.snowflake.hooks.snowflake import SnowflakeHook
2  from airflow.models import Variable
3
4  # Connection Id is read from Airflow Variable, defaults to 'snowflake_catfish'
5  SNOWFLAKE_CONN_ID = Variable.get("SNOWFLAKE_CONN_ID", default_var="
    snowflake_catfish")

```

```

6
7 @task
8 def load_full_refresh():
9     hook = SnowflakeHook(snowflake_conn_id=SNOWFLAKE_CONN_ID)
10    conn = hook.get_conn()
11    cur = conn.cursor()
12
13    cur.execute("BEGIN")
14    # transactional load into RAW schema
15    cur.executemany(insert_sql, rows)
16    cur.execute("COMMIT")

```

- The Variable SNOWFLAKE_CONN_ID allows flexibility: in the UI you can point it to any valid Airflow Connection without changing code.
- The actual credentials and account/warehouse/database details come from the Airflow Connection with ID snowflake_catfish.
- The task wraps operations inside BEGIN/COMMIT ensuring that inserts are atomic. If an exception occurs, a ROLLBACK can be executed for consistency.

Thus, the connection part in your assignment is implemented via SnowflakeHook, pulling configuration securely from Airflow's Admin → Connections.

Quick validation.

After saving the connection, trigger a DAG run and check task logs. If the conn_id is wrong or fields are missing, Airflow will raise a clear error in the task log (e.g., "The conn_id snowflake_catfish isn't defined").

IV. ADMIN → VARIABLES

An example variable (ALPHAVANTAGE_API_KEY) is configured and consumed in code.

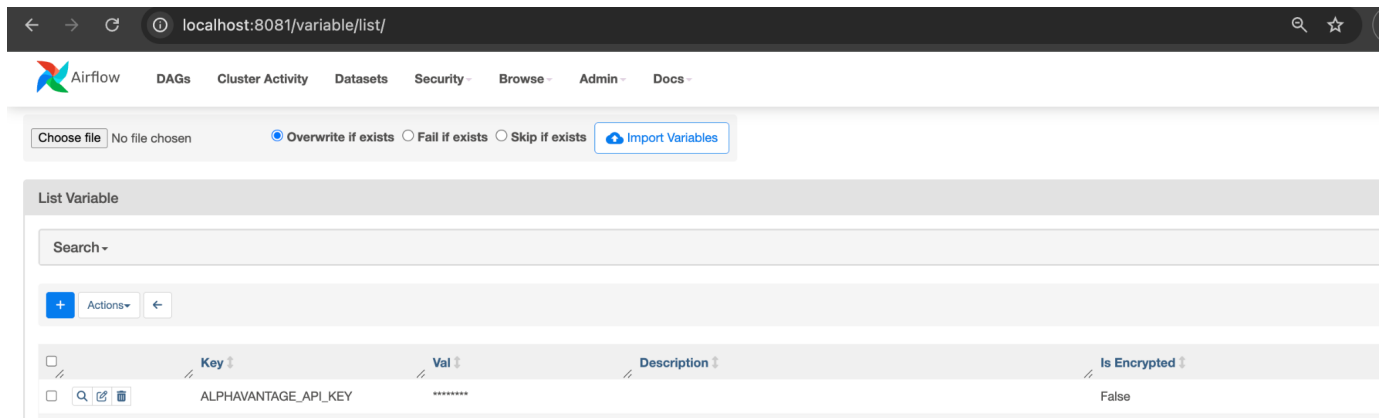


Fig. 3: Airflow Variables page.

How to Add and Use Airflow Variables

Path: Admin → Variables → +

1) Click + and set:

- **Key:** ALPHAVANTAGE_API_KEY
- **Val:** <your-actual-api-key>

2) Save and verify it appears in the list (Figure 4 shows it in use via successful runs).

Where it is used in code (exact snippet).

```
1 from airflow.models import Variable
2
3 API_KEY = Variable.get("ALPHAVANTAGE_API_KEY") # raises if missing
4
5 # Optionally make it non-fatal with a default:
6 # API_KEY = Variable.get("ALPHAVANTAGE_API_KEY", default_var=None)
```

JSON and typed Variables (optional).

When storing structured config (symbols, thresholds, etc.), keep it as JSON and parse it in code:

```
1 import json
2
3 symbols_json = Variable.get("stock_symbols", default_var='["AAPL", "MSFT", "TSLA"]')
4 SYMBOLS = json.loads(symbols_json) # -> ["AAPL", "MSFT", "TSLA"]
5
6 lookback_days = int(Variable.get("lookback_days", default_var="365"))
```

Best practices.

- Use Variables for *non-secret* runtime knobs (symbols, lookback windows, flags).
- Keep *secrets* (passwords, tokens) in *Connections* or a secrets backend (Vault, AWS SM, etc.).
- Validate presence of critical Variables early in the DAG and fail fast with a clear message.

V. RELEVANT CODE SNIPPETS

The following excerpts from `Assignment_5.py` capture the important mechanics: connection retrieval, ETL/ML orchestration, and transactional load into Snowflake. The full file is provided in the project and may be included as an appendix if required.

A. Core DAG Logic

```
1 # dags/assignment_alpha_simple_ctx.py
2 from airflow import DAG
3 from airflow.decorators import task
4 from airflow.models import Variable
5 from airflow.providers.snowflake.hooks.snowflake import SnowflakeHook
6
7 from datetime import datetime
8 import requests
9
10 # Read the Snowflake connection id from a Variable
11 SNOWFLAKE_CONN_ID = Variable.get("SNOWFLAKE_CONN_ID", default_var="
    snowflake_catfish")
12
13 with DAG(
14     dag_id="assignment_alpha_simple_ctx",
15     start_date=datetime(2025, 9, 1),
16     schedule="@daily",
17     catchup=False,
18     tags=["assignment", "alpha", "snowflake"],
19 ) as dag:
20
21     @task
22     def extract() -> list[dict]:
23         """Pull compact daily OHLCV from Alpha Vantage for a small symbol list."""
24         api_key = Variable.get("ALPHAVANTAGE_API_KEY") # must exist
25         symbols_csv = Variable.get("ALPHAVANTAGE_SYMBOLS", default_var="AAPL,MSFT")
26         )
```

```

26     symbols = [s.strip().upper() for s in symbols_csv.split(",") if s.strip()]
27
28     base = "https://www.alphavantage.co/query"
29     rows: list[dict] = []
30
31     for sym in symbols:
32         r = requests.get(
33             base,
34             params={
35                 "function": "TIME_SERIES_DAILY",
36                 "symbol": sym,
37                 "outputsize": "compact",
38                 "datatype": "json",
39                 "apikey": api_key,
40             },
41             timeout=30,
42         )
43         r.raise_for_status()
44         data = r.json()
45         ts = data.get("Time Series (Daily)")
46         if not ts:
47             continue
48         for ds, f in ts.items():
49             try:
50                 rows.append(
51                     {
52                         "SYMBOL": sym,
53                         "PRICE_DATE": ds,
54                         "OPEN": float(f.get("1. open", 0) or 0),
55                         "HIGH": float(f.get("2. high", 0) or 0),
56                         "LOW": float(f.get("3. low", 0) or 0),
57                         "CLOSE": float(f.get("4. close", 0) or 0),
58                         "VOLUME": int(float(f.get("5. volume", 0) or 0)),
59                         "NOTE": None,
60                     }
61                 )
62             except Exception:
63                 # skip malformed row
64                 pass
65
66         # simple sanity filter
67         return [r for r in rows if r["PRICE_DATE"] and r["CLOSE"] and r["CLOSE"] >
68                 0]
69
70 @task
71 def load_full_refresh(rows: list[dict]) -> str:
72     """
73     Full refresh into RAW.STOCK_PRICES_AV:
74     - CREATE SCHEMA/TABLE IF NOT EXISTS
75     - TRUNCATE
76     - INSERT (executemany)
77     - COMMIT / ROLLBACK
78     """
79     schema = Variable.get("SNOWFLAKE_SCHEMA", default_var="RAW")
80     table = Variable.get("SNOWFLAKE_TABLE", default_var="STOCK_PRICES_AV")
81
82     hook = SnowflakeHook(snowflake_conn_id=SNOWFLAKE_CONN_ID)

```

```

82     conn = hook.get_conn()
83     cur  = conn.cursor()
84
85     try:
86         cur.execute("BEGIN")
87         cur.execute(f"CREATE SCHEMA IF NOT EXISTS {schema}")
88         cur.execute(f"""
89             CREATE TABLE IF NOT EXISTS {schema}.{table} (
90                 SYMBOL      STRING,
91                 PRICE_DATE  DATE,
92                 OPEN        FLOAT,
93                 HIGH        FLOAT,
94                 LOW         FLOAT,
95                 CLOSE       FLOAT,
96                 VOLUME      NUMBER,
97                 NOTE        STRING
98             )
99             """)
100        cur.execute(f"TRUNCATE TABLE {schema}.{table}")
101
102        if rows:
103            insert_sql = f"""
104                INSERT INTO {schema}.{table}
105                (SYMBOL, PRICE_DATE, OPEN, HIGH, LOW, CLOSE, VOLUME, NOTE)
106                VALUES (%(SYMBOL)s, %(PRICE_DATE)s, %(OPEN)s, %(HIGH)s, %(LOW)
107                    s, %(CLOSE)s, %(VOLUME)s, %(NOTE)s)
108                """
109            cur.executemany(insert_sql, rows)
110
111            cur.execute("COMMIT")
112            return f"Loaded {len(rows)} rows into {schema}.{table}"
113        except Exception:
114            cur.execute("ROLLBACK")
115            raise
116        finally:
117            cur.close()
118            conn.close()
119
120    # wiring
121    load_full_refresh(extract())

```

B. SQL Transaction for Full Refresh (If Implemented in Python Task)

If the DAG uses Python to run SQL, ensure BEGIN / COMMIT / ROLLBACK are used appropriately, and MERGE/DELETE+INSERT are atomic.

```

1  def load_full_refresh(rows: list[dict]) -> str:
2      """
3      Full refresh into RAW.STOCK_PRICES_AV:
4      - CREATE SCHEMA/TABLE IF NOT EXISTS
5      - TRUNCATE
6      - INSERT (executemany)
7      - COMMIT / ROLLBACK
8      """
9      schema = Variable.get("SNOWFLAKE_SCHEMA", default_var="RAW")
10     table  = Variable.get("SNOWFLAKE_TABLE", default_var="STOCK_PRICES_AV")
11

```

```

12     hook = SnowflakeHook(snowflake_conn_id=SNOWFLAKE_CONN_ID)
13     conn = hook.get_conn()
14     cur = conn.cursor()
15
16     try:
17         cur.execute("BEGIN")
18         cur.execute(f"CREATE SCHEMA IF NOT EXISTS {schema}")
19         cur.execute(f"""
20             CREATE TABLE IF NOT EXISTS {schema}.{table} (
21                 SYMBOL      STRING,
22                 PRICE_DATE  DATE,
23                 OPEN        FLOAT,
24                 HIGH        FLOAT,
25                 LOW         FLOAT,
26                 CLOSE       FLOAT,
27                 VOLUME      NUMBER,
28                 NOTE        STRING
29             )
30         """)
31         cur.execute(f"TRUNCATE TABLE {schema}.{table}")
32
33         if rows:
34             insert_sql = f"""
35                 INSERT INTO {schema}.{table}
36                 (SYMBOL, PRICE_DATE, OPEN, HIGH, LOW, CLOSE, VOLUME, NOTE)
37                 VALUES (%(SYMBOL)s, %(PRICE_DATE)s, %(OPEN)s, %(HIGH)s, %(LOW)
38                     s, %(CLOSE)s, %(VOLUME)s, %(NOTE)s)
39             """
40             cur.executemany(insert_sql, rows)
41
42             cur.execute("COMMIT")
43             return f"Loaded {len(rows)} rows into {schema}.{table}"
44         except Exception:
45             cur.execute("ROLLBACK")
46             raise
47         finally:
48             cur.close()
49             conn.close()
50
51     # wiring
52     load_full_refresh(extract())

```

VI. EXECUTION EVIDENCE: EVENT LOGS

Figures 4 and 5 show successful runs of the two primary tasks (extract and load_full_refresh).

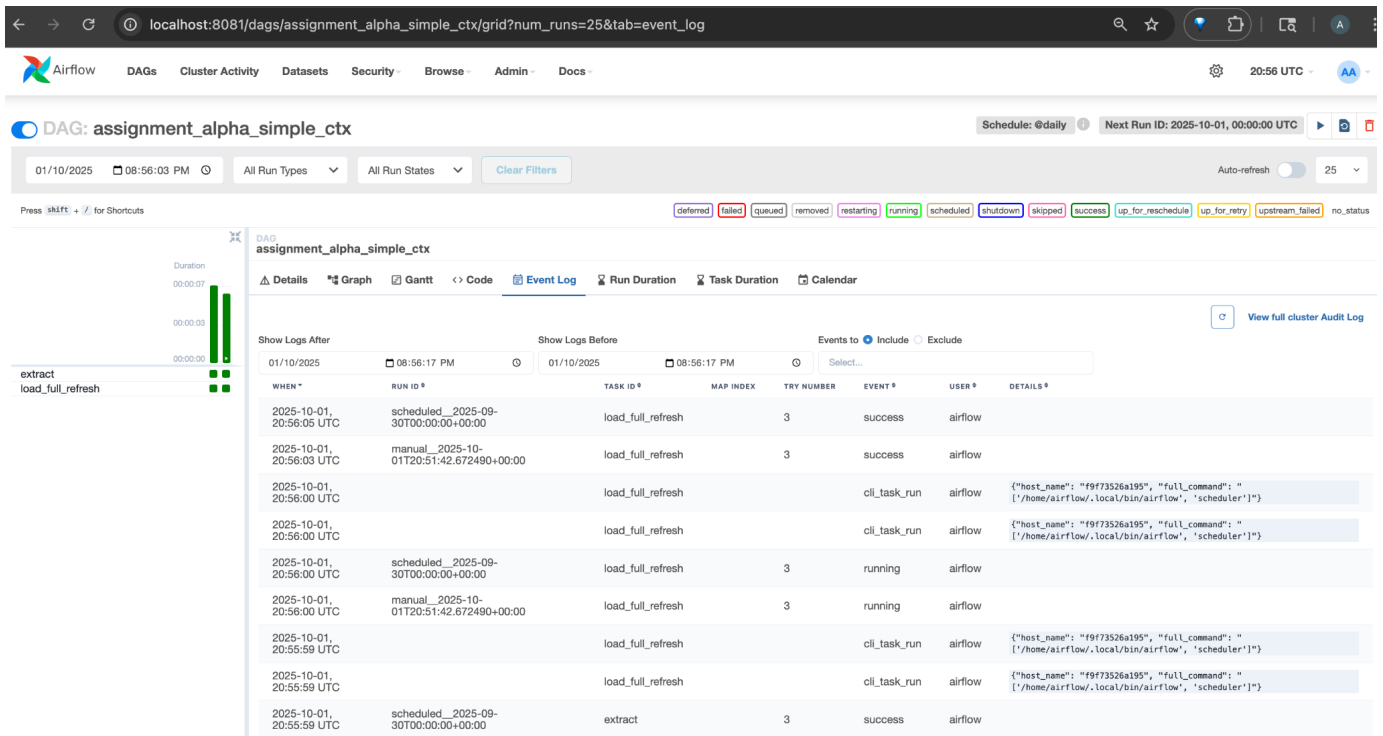


Fig. 4: Event log view with task successes.

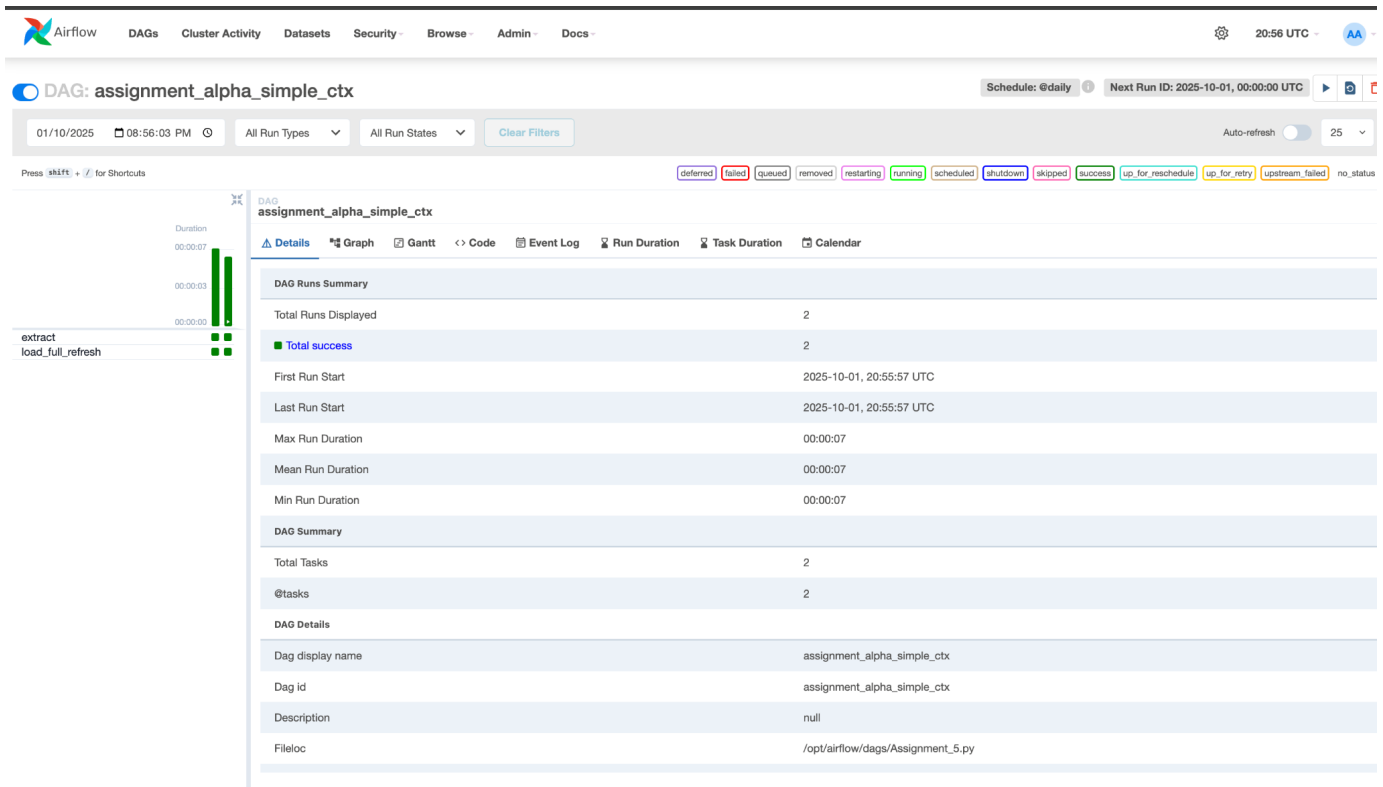


Fig. 5: DAG details summary confirming successful runs.

VII. CONCLUSION

This report demonstrates a working Airflow setup with Snowflake integration, proper configuration via Connections/Variables, transactional loading, and verified successful DAG runs.

APPENDIX (OPTIONAL): FULL CODE

```

1  # dags/assignment_alpha_simple_ctx.py
2  from airflow import DAG
3  from airflow.decorators import task
4  from airflow.models import Variable
5  from airflow.providers.snowflake.hooks.snowflake import SnowflakeHook
6
7  from datetime import datetime
8  import requests
9
10 # Read the Snowflake connection id from a Variable
11 SNOWFLAKE_CONN_ID = Variable.get("SNOWFLAKE_CONN_ID", default_var="
    snowflake_catfish")
12
13 with DAG(
14     dag_id="assignment_alpha_simple_ctx",
15     start_date=datetime(2025, 9, 1),
16     schedule="@daily",
17     catchup=False,
18     tags=["assignment", "alpha", "snowflake"],
19 ) as dag:
20
21     @task
22     def extract() -> list[dict]:
23         """Pull compact daily OHLCV from Alpha Vantage for a small symbol list."""
24         api_key = Variable.get("ALPHAVANTAGE_API_KEY") # must exist
25         symbols_csv = Variable.get("ALPHAVANTAGE_SYMBOLS", default_var="AAPL,MSFT")
26         symbols = [s.strip().upper() for s in symbols_csv.split(",") if s.strip()]
27
28         base = "https://www.alphavantage.co/query"
29         rows: list[dict] = []
30
31         for sym in symbols:
32             r = requests.get(
33                 base,
34                 params={
35                     "function": "TIME_SERIES_DAILY",
36                     "symbol": sym,
37                     "outputsize": "compact",
38                     "datatype": "json",
39                     "apikey": api_key,
40                 },
41                 timeout=30,
42             )
43             r.raise_for_status()
44             data = r.json()
45             ts = data.get("Time Series (Daily)")
46             if not ts:
47                 continue
48             for ds, f in ts.items():
49                 try:
50                     rows.append(
51                         {
52                             "SYMBOL": sym,
53                             "PRICE_DATE": ds,
54                             "OPEN": float(f.get("1. open", 0) or 0),

```

```

55         "HIGH": float(f.get("2. high", 0) or 0),
56         "LOW": float(f.get("3. low", 0) or 0),
57         "CLOSE": float(f.get("4. close", 0) or 0),
58         "VOLUME": int(float(f.get("5. volume", 0) or 0)),
59         "NOTE": None,
60     }
61 )
62 except Exception:
63     # skip malformed row
64     pass
65
66 # simple sanity filter
67 return [r for r in rows if r["PRICE_DATE"] and r["CLOSE"] and r["CLOSE"] >
68         0]
69
70 @task
71 def load_full_refresh(rows: list[dict]) -> str:
72     """
73     Full refresh into RAW.STOCK_PRICES_AV:
74     - CREATE SCHEMA/TABLE IF NOT EXISTS
75     - TRUNCATE
76     - INSERT (executemany)
77     - COMMIT / ROLLBACK
78     """
79     schema = Variable.get("SNOWFLAKE_SCHEMA", default_var="RAW")
80     table = Variable.get("SNOWFLAKE_TABLE", default_var="STOCK_PRICES_AV")
81
82     hook = SnowflakeHook(snowflake_conn_id=SNOWFLAKE_CONN_ID)
83     conn = hook.get_conn()
84     cur = conn.cursor()
85
86     try:
87         cur.execute("BEGIN")
88         cur.execute(f"CREATE SCHEMA IF NOT EXISTS {schema}")
89         cur.execute(f"""
90             CREATE TABLE IF NOT EXISTS {schema}.{table} (
91                 SYMBOL      STRING,
92                 PRICE_DATE DATE,
93                 OPEN         FLOAT,
94                 HIGH         FLOAT,
95                 LOW          FLOAT,
96                 CLOSE        FLOAT,
97                 VOLUME       NUMBER,
98                 NOTE         STRING
99             )
100         """)
101         cur.execute(f"TRUNCATE TABLE {schema}.{table}")
102
103         if rows:
104             insert_sql = f"""
105                 INSERT INTO {schema}.{table}
106                 (SYMBOL, PRICE_DATE, OPEN, HIGH, LOW, CLOSE, VOLUME, NOTE)
107                 VALUES (%(SYMBOL)s, %(PRICE_DATE)s, %(OPEN)s, %(HIGH)s, %(LOW)
108                     s, %(CLOSE)s, %(VOLUME)s, %(NOTE)s)
109             """
110             cur.executemany(insert_sql, rows)

```

```
110         cur.execute("COMMIT")
111         return f"Loaded {len(rows)} rows into {schema}.{table}"
112     except Exception:
113         cur.execute("ROLLBACK")
114         raise
115     finally:
116         cur.close()
117         conn.close()
118
119 # wiring
120 load_full_refresh(extract())
```